

Multiclass Support Vector Machine exercise

Complete and hand in the completed notebook (including the output) with your assignment submission. You will be submitting the homework as a zip file including all the parts on the Blackboard.

In this exercise you will:

- implement a full-vectorized loss function for the SVM
- implement the fully-vectorized expression for its analytic gradient
- check your implementation using numerical gradient
- use a validation set to tune the learning rate and regularization strength
- optimize the loss function with SGD
- visualize the final learned weights

```
In [1]: ## Default modules
from random import print, function
import matplotlib.pyplot as plt
import numpy as np
import matplotlib.pyplot as plt
from io import StringIO
from io import StringIO

# This is a bit of magic to make matplotlib figures appear inline in the
# notebook rather than in a new window.
%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# Some more magic so that the notebook will reload external python modules;
# see http://stackoverflow.com/question/17529682
%load_ext autoreload
%autoreload 2
```

CI-FAR-10 Data Loading and Preprocessing

```
In [2]: # Load the raw CIFAR-10 data.
cifar10_dir = 'data/cifar10'
X_train, y_train, X_test, y_test = load_cifar10(cifar10_dir)

# Sanity check, we verify the sizes of the training and test data.
print('Training data shape: ', X_train.shape)
print('Training labels shape: ', y_train.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

In [3]: # Visualize some examples from the dataset.

We display a few examples of training images from each class.

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

num_classes = len(classes)

samples_per_class = 7

for y, cls in enumerate(classes):

 idxs = np.flatnonzero(y_train == y)

 for i, idx in enumerate(idxs):

 plt_idx = i % num_classes + y + 1

 plt.imshow(X_train[idxs.astype('int64')])

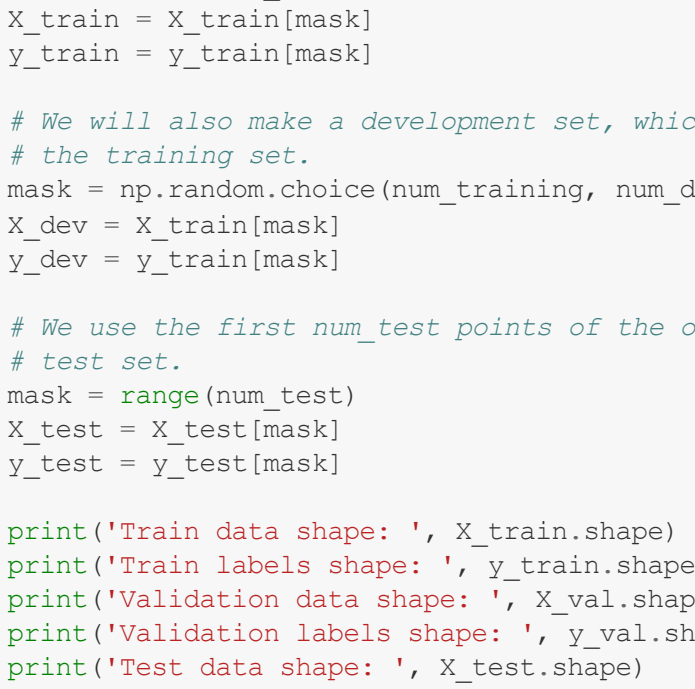
 plt.axis('off')

 if i % 10 == 0:

 plt.title(cls)

plt.show()

plane car bird cat deer dog frog horse ship truck



```
In [4]: # Split the data into train, val, and test sets. In addition we will
# create a small development set as a subset of the training data;
# we can use this 2% development set to tune the model.
num_training = 49000
num_validation = 1000
num_test = 1000
num_dev = 500

# Our validation set will be num_validation points from the original
# training set.
mask = range(num_training, num_training + num_validation)
X_val = X_train[mask]
y_val = y_train[mask]

# Our training set will be the first num_train points from the original
# training set.
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

# We will also make a development set, which is a small subset of
# the training set.
mask = np.random.choice(num_training, num_dev, replace=False)
X_dev = X_train[mask]
y_dev = y_train[mask]
```

```
# We use the first num_test points of the original test set as our
# test set.
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)
```

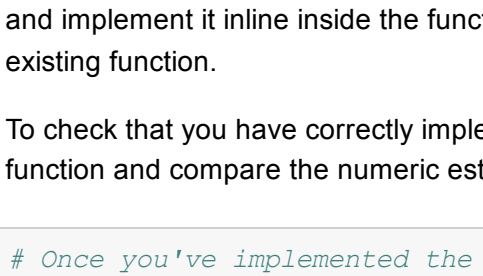
```
In [5]: # Preprocessing: reshape the image data into rows
X_train = np.reshape(X_train, (X_train.shape[0], -1))
X_val = np.reshape(X_val, (X_val.shape[0], -1))
X_test = np.reshape(X_test, (X_test.shape[0], -1))
X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))

# As a sanity check, print out the shapes of the data
print('Training data shape: ', X_train.shape)
print('Validation data shape: ', X_val.shape)
print('Test data shape: ', X_test.shape)
print('Dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

```
In [6]: # Preprocessing: subtract the mean image
# first: compute the image mean based on the training data
mean_image = np.mean(X_train, axis=0)
print('mean image: %s' % str(mean_image))
plt.figure(figsize=(4, 4))
plt.imshow(mean_image, reshape=(32, 32, 3).astype('uint8')) # visualize the mean image
plt.show()
```

mean_image = [130.64189796 135.98173469 132.47391837 130.05569388 135.34804082 131.75402041 130.96055102 136.14328571 132.47636735 131.14467347]



```
In [7]: # second: subtract the mean image from train and test data
X_train -= mean_image
X_val -= mean_image
X_test -= mean_image
X_dev -= mean_image
```

```
In [8]: # third: append the bias dimension of ones (i.e. bias trick) so that our SVM
# only has to worry about optimizing a single weight matrix W
X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])

print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)

(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)
```

SVM Classifier

Your code for this section will all be written inside `le900/classifier/linear_svm.py`.

As you can see, we have pre-filled the function `compute_loss_naive` which uses for loops to evaluate the multiclass SVM loss function.

```
In [9]: # Evaluate the naive implementation of the loss we provided for you:
from le900.classifier.linear_svm import svm_loss_naive
import time

# generate a random SVM weight matrix of small numbers
W = np.random.randn(3073, 10) * 0.0001

loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.000004)
print('loss: %f' % (loss,))

loss: 9.223893
```

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have correctly implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you:

```
In [10]: # Once you've implemented the gradient, recompute it with the code below
# and gradient check it with the function we provided for you

# Compute the loss and its gradient at W.
loss, grad = svm_loss_naive(W, X_dev, y_dev, 0.0)

# Numerically compute the gradient along several randomly chosen dimensions, and
# compare them with your analytically computed gradient. The numbers should match
# almost exactly.
from le900.gradient_check import grad_check_sparse
f = lambda w: svm_loss_naive(w, X_dev, y_dev, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad)

numerical: 7.947813 analytic: 7.947813, relative error: 2.175078e-11
numerical: -9.409788 analytic: -9.409788, relative error: 6.969942e-11
numerical: -10.521413 analytic: -10.521413, relative error: 2.128826e-11
numerical: 9.810162 analytic: 9.810162, relative error: 6.49475e-12
numerical: 14.669548 analytic: 14.669548, relative error: 1.455233e-11
numerical: -6.361561 analytic: -6.361561, relative error: 3.288525e-12
numerical: 12.129248 analytic: 12.129248, relative error: 1.614808e-11
numerical: -55.304365 analytic: -55.304365, relative error: 5.833571e-13
numerical: -0.828118 analytic: -0.828118, relative error: 4.265895e-10
numerical: 2.122222e-07 req 2.122222e-07 train accuracy: 0.34566 val accuracy: 0.34566
numerical: -16.870974 analytic: -16.870974, relative error: 2.694021e-11
numerical: -1.226697 analytic: -1.226697, relative error: 3.368506e-10
numerical: -0.618691 analytic: -0.618691, relative error: 3.454566e-11
numerical: -0.796062 analytic: -0.796062, relative error: 5.749047e-10
numerical: -14.698456 analytic: -14.698456, relative error: 2.448776e-12
numerical: -5.343436 analytic: -5.343436, relative error: 2.253938e-11
numerical: -60.694161 analytic: -60.694161, relative error: 4.395176e-11
numerical: 3.224499 analytic: 3.224499, relative error: 2.530766e-11
numerical: 9.688917 analytic: 9.688917, relative error: 2.017144e-11
```

Inline Question 1:

It is possible that once in a while a dimension in the gradcheck will not match exactly. What could such a discrepancy be caused by? It is a reason for concern? What is a simple example in one dimension where a gradient check could fail? Hint: the SVM loss function is not strictly speaking differentiable

Your Answer: One reason could be the discontinuity of SVM loss function when $S_j - f_j = 1$ for $j = 0, 1, \dots, 9$, which leads to an ill-defined gradient at that point. Numerically, when the function hits that point, the numerical error could be very large, thereby leads to the possible fail of gradient check.

```
In [11]: # Next implement the function svm_loss_vectorized; for now only compute the loss;
# we will implement the gradient in a moment.
tic = time.time()
loss_naive, grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000004)
toc = time.time()
print('Naive loss: %f computed in %fs' % (loss_naive, toc - tic))

from le900.classifier.linear_svm import svm_loss_vectorized
tic = time.time()
loss_vectorized, grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000004)
toc = time.time()
print('Vectorized loss: %f computed in %fs' % (loss_vectorized, toc - tic))

# The losses should match but your vectorized implementation should be much faster.
print('Difference: %f' % (loss_naive - loss_vectorized))

Naive loss: 9.223893e+00 computed in 0.08512s
Vectorized loss: 9.223893e+00 computed in 0.006516s
Difference: 0.000000
```

```
In [12]: # Complete the implementation of svm_loss_vectorized, and compute the gradient
# of the loss function in a vectorized way.

# The naive implementation and the vectorized implementation should match, but
# the vectorized version should still be much faster.
tic = time.time()
grad_naive = svm_loss_naive(W, X_dev, y_dev, 0.000004)
toc = time.time()
print('Naive loss and gradient: computed in %fs' % (toc - tic))

tic = time.time()
grad_vectorized = svm_loss_vectorized(W, X_dev, y_dev, 0.000004)
toc = time.time()
print('Vectorized loss and gradient: computed in %fs' % (toc - tic))

# The loss is a single number, so it is easy to compare the values computed
# by the two implementations. The gradient on the other hand is a matrix, so
# we use the Frobenius norm to compare them.
difference = np.linalg.norm(grad_naive - grad_vectorized, ord='fro')
print('Difference: %f' % difference)

Naive loss and gradient: computed in 0.085052s
Vectorized loss and gradient: computed in 0.006338s
Difference: 0.000000
```

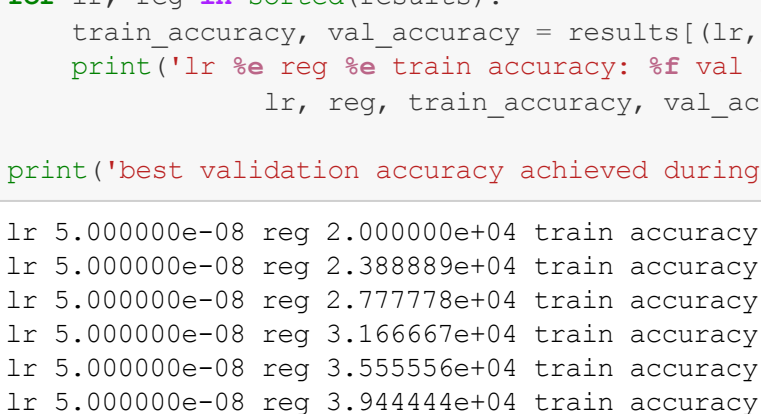
Stochastic Gradient Descent

We now have vectorized and efficient expressions for the loss, the gradient and our gradient matches the numerical gradient. We are therefore ready to do SGD to minimize the loss.

```
In [13]: # In the file linear_classifier.py, implement SGD in the function
# linear_classifier_train() and then run it with the code below.
from le900.classifier import linear_svm
svm = LinearSVM()
tic = time.time()
loss_hist = svm.train(X_train, y_train, learning_rate=1e-4, num_iters=100, verbose=True)
toc = time.time()
print('That took %fs' % (toc - tic))

iteration 0 / 1000: loss 44.091611
iteration 100 / 1000: loss 193.641019
iteration 200 / 1000: loss 61.067365
iteration 300 / 1000: loss 21.638116
iteration 400 / 1000: loss 9.847789
iteration 500 / 1000: loss 7.007161
iteration 600 / 1000: loss 5.662183
iteration 700 / 1000: loss 5.467002
iteration 800 / 1000: loss 5.700516
iteration 900 / 1000: loss 5.217007
iteration 1000 / 1000: loss 4.659286
iteration 1100 / 1000: loss 5.372933
iteration 1200 / 1000: loss 5.225398
iteration 1300 / 1000: loss 5.707822
iteration 1400 / 1000: loss 5.591786
That took 6.043468s
```

```
In [14]: # A useful debugging strategy is to plot the loss as a function of
# iteration number.
plt.plot(loss_hist)
plt.xlabel('Iteration number')
plt.ylabel('Loss value')
plt.show()
```



```
In [15]: # Write the linearSVM.predict function and evaluate the performance on both the
# training and validation set
y_train_pred = svm.predict(X_train)
print('training accuracy: %f' % (np.mean(y_train == y_train_pred),))
y_val_pred = svm.predict(X_val)
print('validation accuracy: %f' % (np.mean(y_val == y_val_pred),))

training accuracy: 0.37013
validation accuracy: 0.39000
```

```
In [29]: # Use the validation set to tune hyperparameters (regularization strength and
# learning rate). You should experiment with different ranges for the learning
# rate and regularization strength; if you are careful, you should be able to
# get a classification accuracy of about 0.4 on the validation set.
learning_rates = [1e-5, 5e-5]
regularization_strengths = [2e4, 5e4]
```

```
# results is dictionary mapping tuples of the form
# (learning_rate, regularization_strength) to tuples of the form
# (training_accuracy, validation_accuracy). The accuracy is simply the fraction
# of data points that are correctly classified.
results = {}
best_val = -1 # The highest validation accuracy that we have seen so far.
best_svm = None # The LinearSVM object that achieved the highest validation rate.
```

```
#####
# TODO:
# Write code that chooses the best hyperparameters by tuning on the validation set.
# For each combination of hyperparameters, train a Linear SVM on the
# training set, compute its accuracy on the training and validation sets, and
# store these numbers in the results dictionary. In addition, store the best
# validation accuracy in best_val and the LinearSVM object that achieves this
# accuracy in best_svm.
# Hint: you should use a small value for num_iters as you develop your
# validation code so that the SVMs don't take much time to train; once you are
# confident that your validation code works, you should return the validation
# code with a larger value for num_iters.
#####
START OF YOUR CODE
pass # Write your code here
```

```
# Define new lr and reg ranges
learning_rates = [1e-5, 5e-5]
regularization_strengths = [2e4, 5e4]
for lr in np.linspace(learning_rates[0], learning_rates[1], num=10):
    for reg in np.linspace(regularization_strengths[0], regularization_strengths[1], num=10):
        svm = LinearSVM()
        train_acc = svm.predict(X_train)
        val_acc = np.mean(y_val == y_val_pred)
        val_accuracy = np.mean(y_val == y_val_pred)
```

```
        results[(lr, reg)] = (train_accuracy, val_accuracy)
        if val_accuracy > best_val:
            best_val = val_accuracy
            best_svm = svm
```

```
#####
END OF YOUR CODE
# Print out results.
for lr, reg in sorted(results):
    train_accuracy, val_accuracy = results[(lr, reg)]
    print('lr: %f reg: %f train accuracy: %f val accuracy: %f' % (
        lr, reg, train_accuracy, val_accuracy))
```

```
print('best validation accuracy achieved during cross-validation: %f' % best_val)

lr: 1e-05 reg: 2.000000e+04 train accuracy: 0.325796 val accuracy: 0.319000
lr: 5.000000e-05 reg: 2.388889e+04 train accuracy: 0.340082 val accuracy: 0.380000
lr: 0.000100 reg: 2.777778e+04 train accuracy: 0.350510 val accuracy: 0.370000
lr: 0.000150 reg: 3.166667e+04 train accuracy: 0.348587 val accuracy: 0.360000
lr: 0.000200 reg: 3.555556e+04 train accuracy: 0.348122 val accuracy: 0.360000
lr: 0.000250 reg: 3.944444e+04 train accuracy: 0.359388 val accuracy: 0.368000
lr: 0.000300 reg: 4.333333e+04 train accuracy: 0.369331 val accuracy: 0.368000
lr: 0.000350 reg: 4.722222e+04 train accuracy: 0.359490 val accuracy: 0.368000
lr: 0.000400 reg: 5.111111e+04 train accuracy: 0.359490 val accuracy: 0.368000
lr: 0.000450 reg: 5.500000e+04 train accuracy: 0.368918 val accuracy: 0.363000
lr: 0.000500 reg: 5.888889e+04 train accuracy: 0.363796 val accuracy: 0.370000
lr: 0.000550 reg: 6.277778e+04 train accuracy: 0.369755 val accuracy: 0.383000
lr: 0.000600 reg: 6.666667e+04 train accuracy: 0.368593 val accuracy: 0.362000
lr: 0.000650 reg: 7.055556e+04 train accuracy: 0.363796 val accuracy: 0.370000
lr: 0.000700 reg: 7.444444e+04 train accuracy: 0.363796 val accuracy: 0.362000
lr: 0.000750 reg: 7.833333e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.000800 reg: 8.222222e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.000850 reg: 8.611111e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.000900 reg: 9.000000e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.000950 reg: 9.388889e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001000 reg: 9.777778e+04 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001050 reg: 1.016667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001100 reg: 1.055556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001150 reg: 1.094444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001200 reg: 1.133333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001250 reg: 1.172222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001300 reg: 1.211111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001350 reg: 1.250000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001400 reg: 1.288889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001450 reg: 1.327778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001500 reg: 1.366667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001550 reg: 1.405556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001600 reg: 1.444444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001650 reg: 1.483333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001700 reg: 1.522222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001750 reg: 1.561111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001800 reg: 1.600000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001850 reg: 1.638889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001900 reg: 1.677778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.001950 reg: 1.716667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002000 reg: 1.755556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002050 reg: 1.794444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002100 reg: 1.833333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002150 reg: 1.872222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002200 reg: 1.911111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002250 reg: 1.950000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002300 reg: 1.988889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002350 reg: 2.027778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002400 reg: 2.066667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002450 reg: 2.105556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002500 reg: 2.144444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002550 reg: 2.183333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002600 reg: 2.222222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002650 reg: 2.261111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002700 reg: 2.300000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002750 reg: 2.338889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002800 reg: 2.377778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002850 reg: 2.416667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002900 reg: 2.455556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.002950 reg: 2.494444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003000 reg: 2.533333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003050 reg: 2.572222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003100 reg: 2.611111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003150 reg: 2.650000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003200 reg: 2.688889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003250 reg: 2.727778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003300 reg: 2.766667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003350 reg: 2.805556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003400 reg: 2.844444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003450 reg: 2.883333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003500 reg: 2.922222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003550 reg: 2.961111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003600 reg: 3.000000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003650 reg: 3.038889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003700 reg: 3.077778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003750 reg: 3.116667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003800 reg: 3.155556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003850 reg: 3.194444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003900 reg: 3.233333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.003950 reg: 3.272222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004000 reg: 3.311111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004050 reg: 3.350000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004100 reg: 3.388889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004150 reg: 3.427778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004200 reg: 3.466667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004250 reg: 3.505556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004300 reg: 3.544444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004350 reg: 3.583333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004400 reg: 3.622222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004450 reg: 3.661111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004500 reg: 3.700000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004550 reg: 3.738889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004600 reg: 3.777778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004650 reg: 3.816667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004700 reg: 3.855556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004750 reg: 3.894444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004800 reg: 3.933333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004850 reg: 3.972222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004900 reg: 4.011111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.004950 reg: 4.050000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005000 reg: 4.088889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005050 reg: 4.127778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005100 reg: 4.166667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005150 reg: 4.205556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005200 reg: 4.244444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005250 reg: 4.283333e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005300 reg: 4.322222e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005350 reg: 4.361111e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005400 reg: 4.400000e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005450 reg: 4.438889e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005500 reg: 4.477778e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005550 reg: 4.516667e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005600 reg: 4.555556e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005650 reg: 4.594444e+05 train accuracy: 0.355857 val accuracy: 0.370000
lr: 0.005700
```